



# Multidatabase Distributed RDF Graph Management

---

Peng Peng  
Hunan University, China  
[hnu16pp@hnu.edu.cn](mailto:hnu16pp@hnu.edu.cn)

# Outline

1. Background
2. Partitioning-agnostic Approaches
3. Federated Systems
4. Conclusions

# **Part 1**

## Background

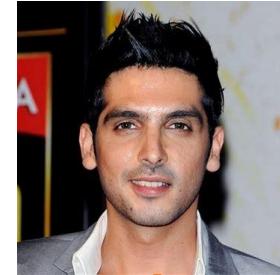
# RDF Introduction

---

## Resource Description Framework

- It is a data model proposed widely used for knowledge bases
- Everything is an uniquely named resource
- Properties of resources can be defined
- Relationships with other resources can be defined

dbpedia:Zayed\_Khan



dbpedia:name "Zayed Khan"@en  
dbpedia:dateOfBirth "1980-07-05"

dbpedia:birthPlace



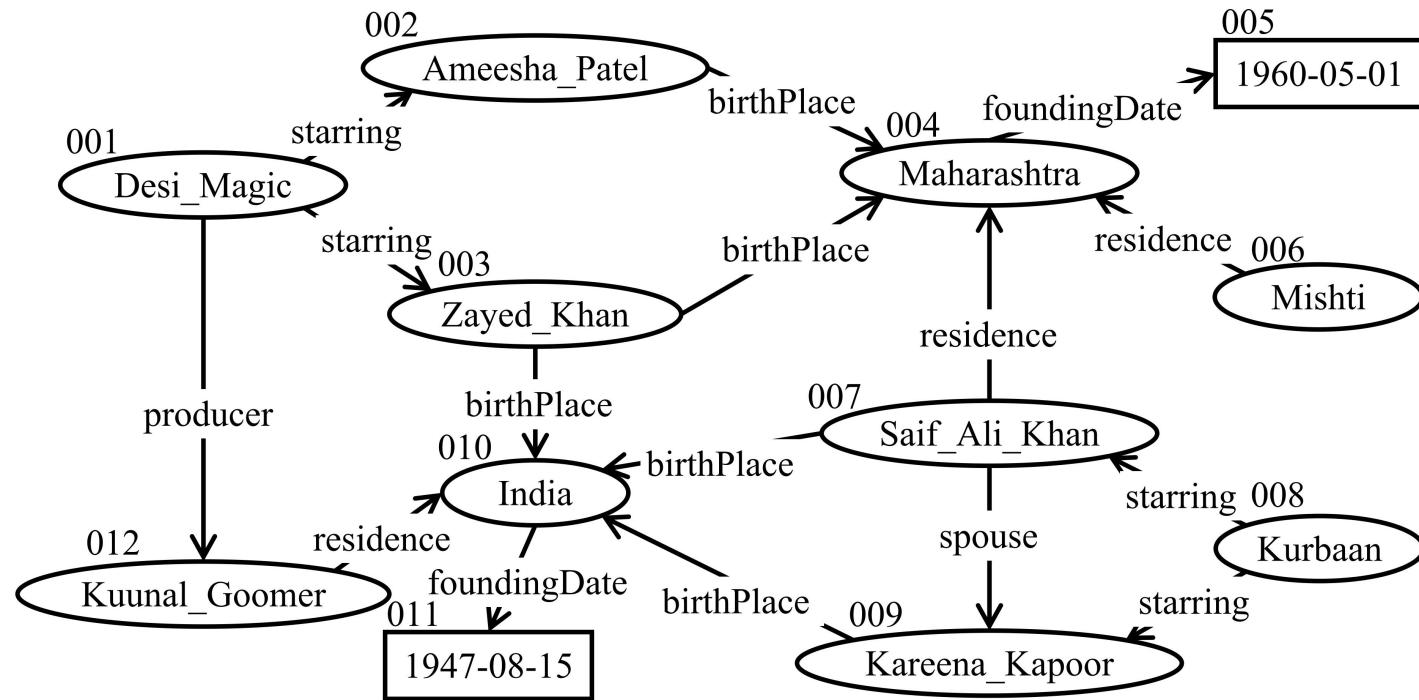
Maharashtra

dbpedia:Maharashtra

# RDF Graph

---

An RDF dataset can be represented as a graph



# RDF Query Model

---

Query Model - SPARQL Protocol and RD<sup>F</sup> Query Language

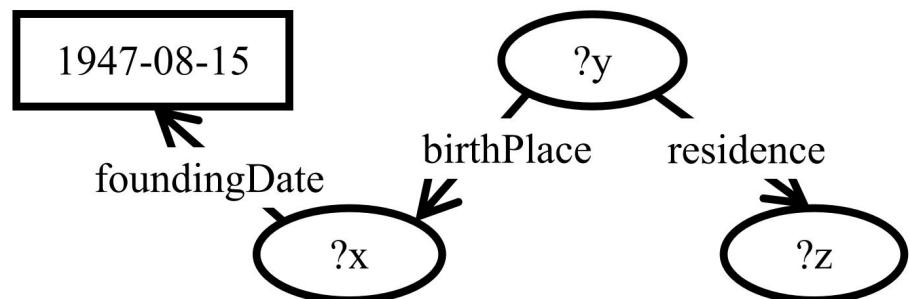
A SPARQL query is a set of triple patterns with variables

Select ?x where {

?y residence ?z .

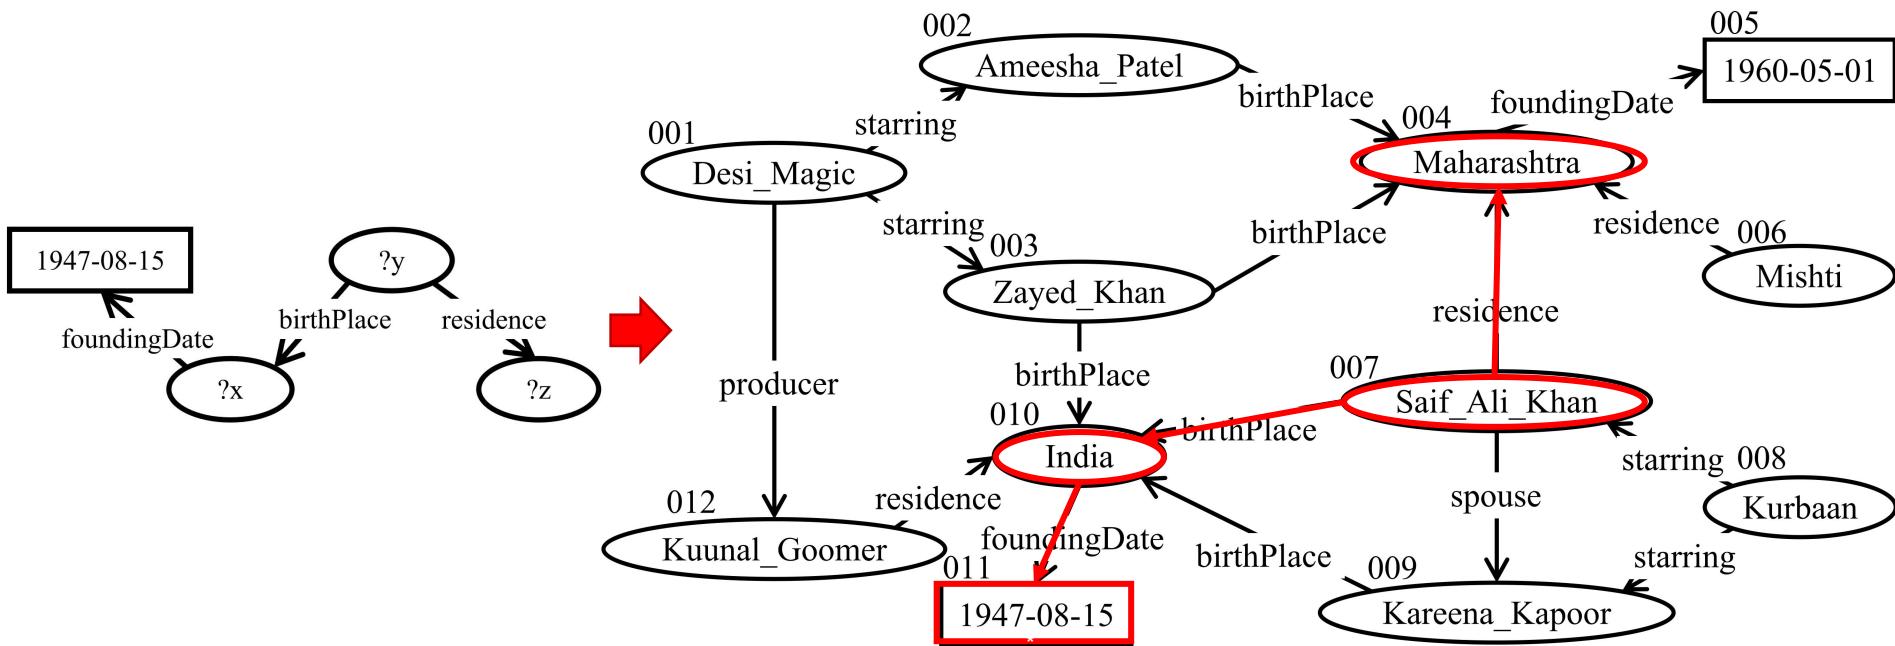
?y birthPlace ?x .

?x foundingDate 1947-08-15 . }



# RDF Query Model

Answering SPARQL query = subgraph matching using homomorphism



# Large RDF Graphs

---

Now, RDF datasets become larger and larger



**Max-Planck-Institute**

284 million triples



**Metaweb Company  
acquired by Google in 2010**

2.4 billion triples



**Leipzig University  
University of Mannheim  
OpenLink Software**

9.5 billion triples

It is important to design a distributed RDF system to manage the large RDF dataset.

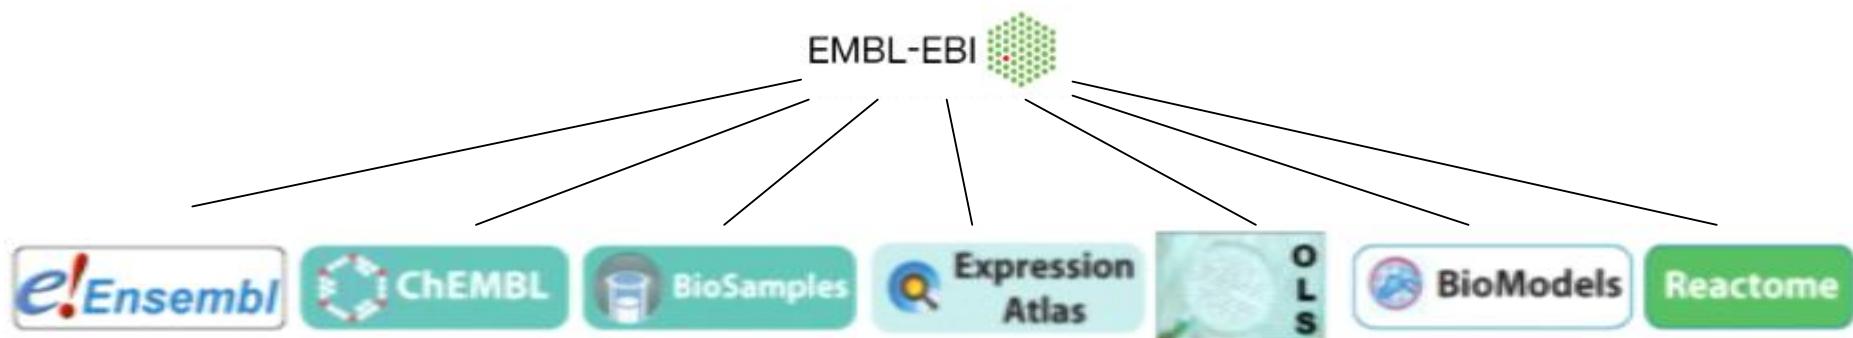
# **Part 2**

Partitioning-agnostic Approaches

# Partitioning-agnostic Application

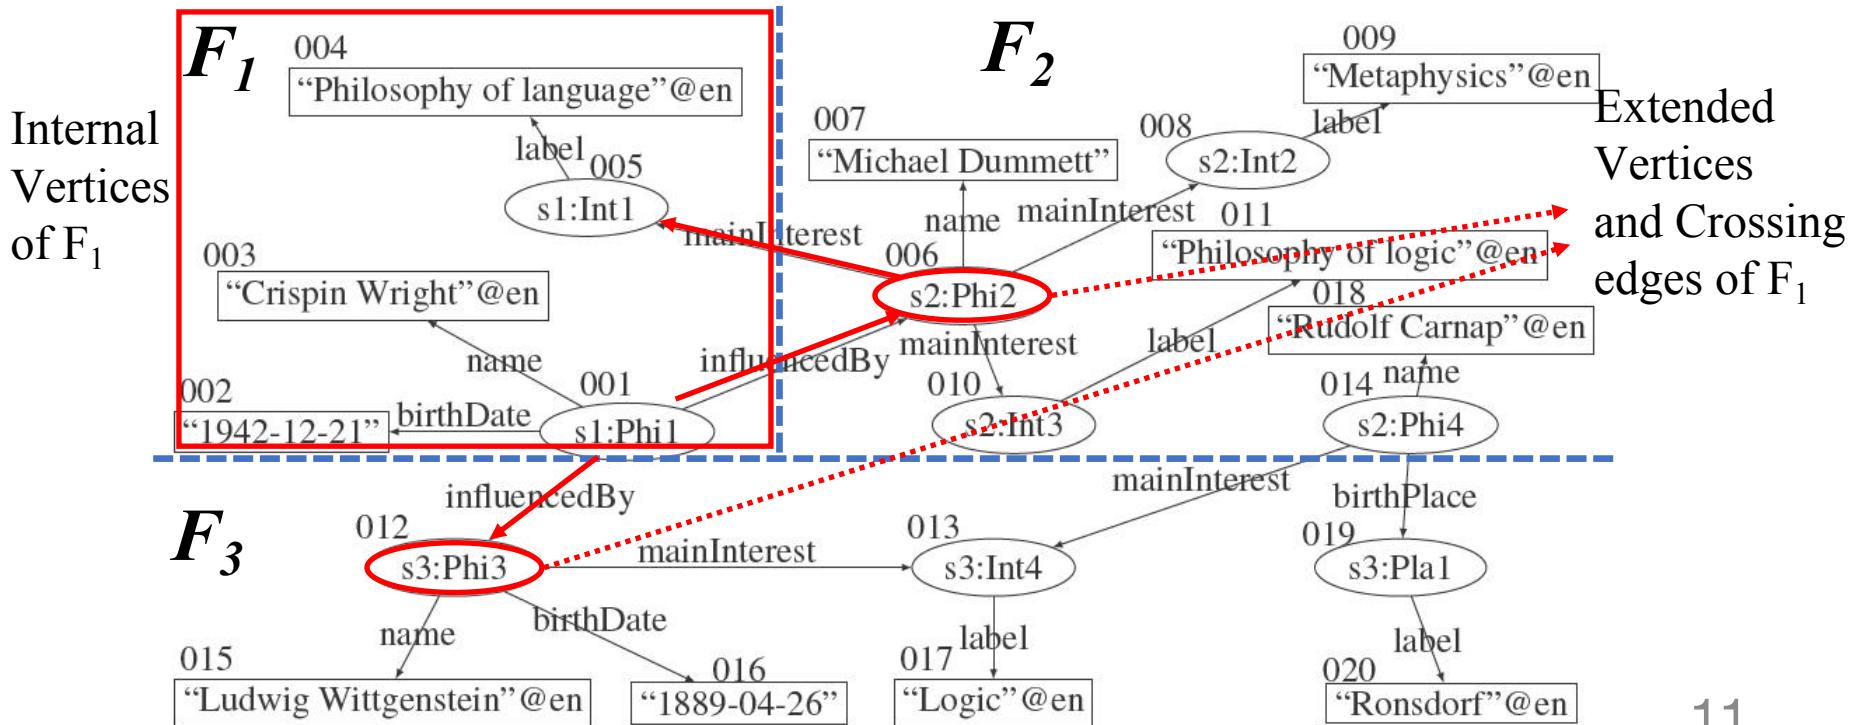
---

- ❑ In some applications, the RDF graph partitioning strategy is not controlled by the distributed RDF system itself



# Distributed RDF Graphs

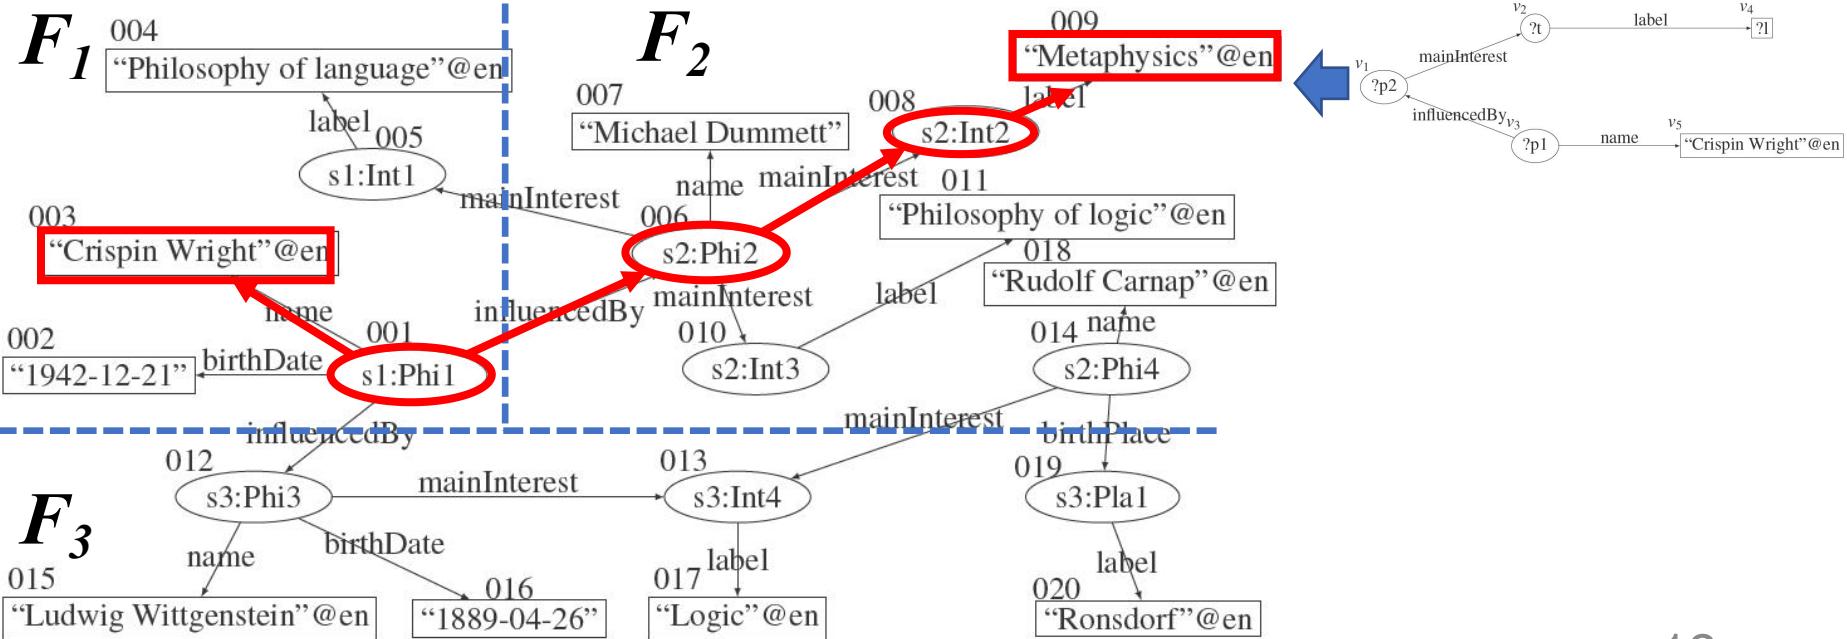
- An RDF graph is vertex-disjoint partitioned through different sites



# Partial Evaluation

- “Partial evaluation” framework is used for distributed SPARQL

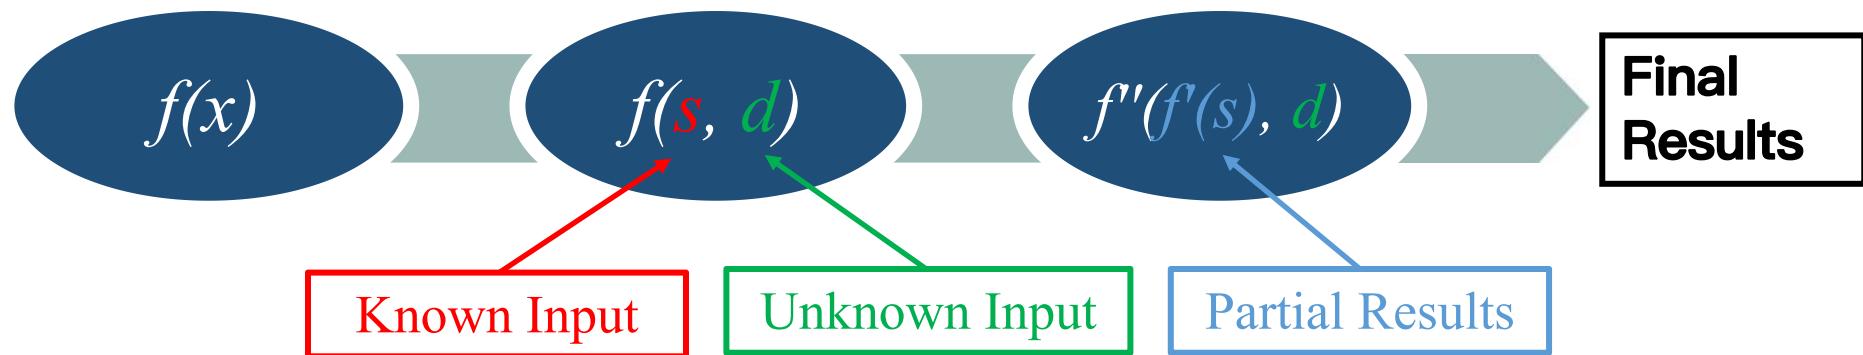
query evaluation



# Partial Evaluation

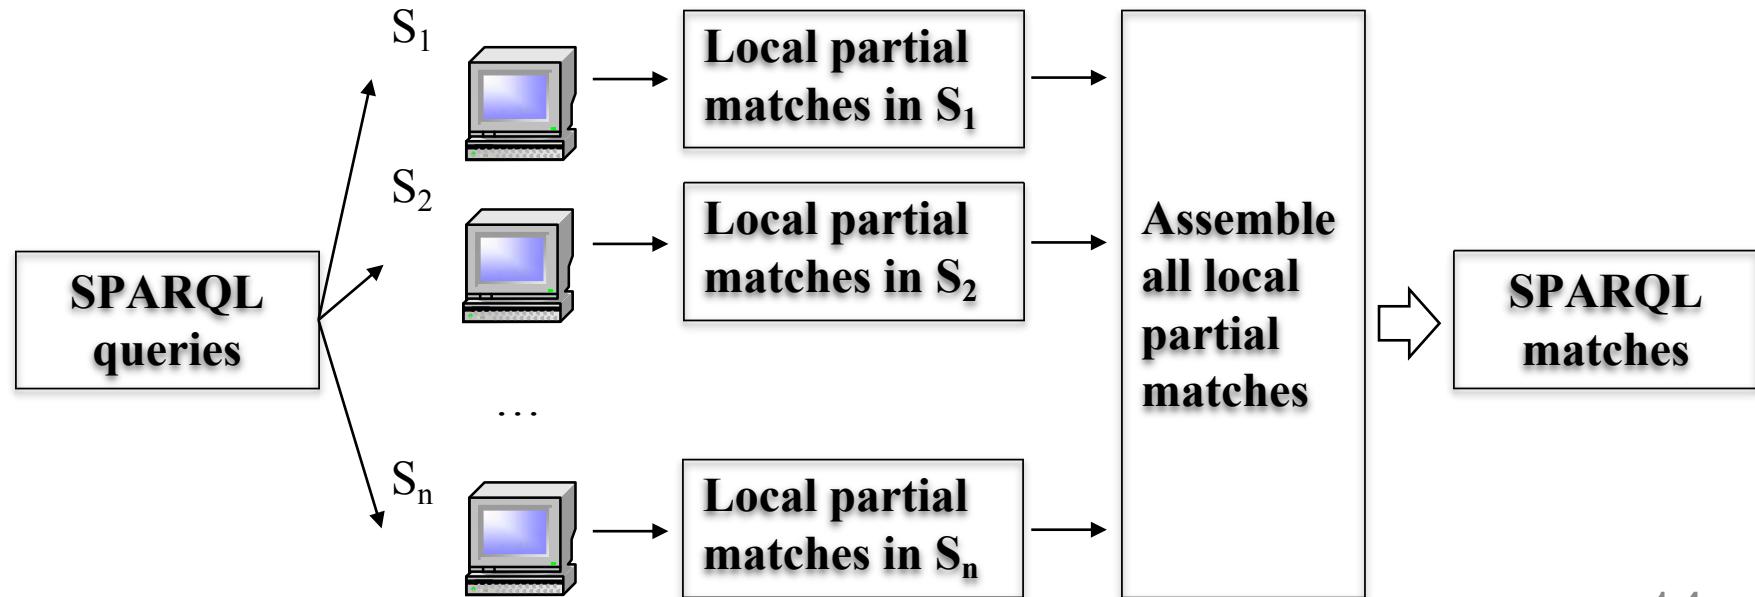
---

- Partial Evaluation is a framework widely used for evaluating queries on distributed graphs



# Partial Evaluation Framework [Peng et al., VLDBJ 2016, ICDE 2019]

- We propose a distributed RDF system, named gStore<sup>D</sup>, to evaluate SPARQL queries in a “partial evaluation” framework.



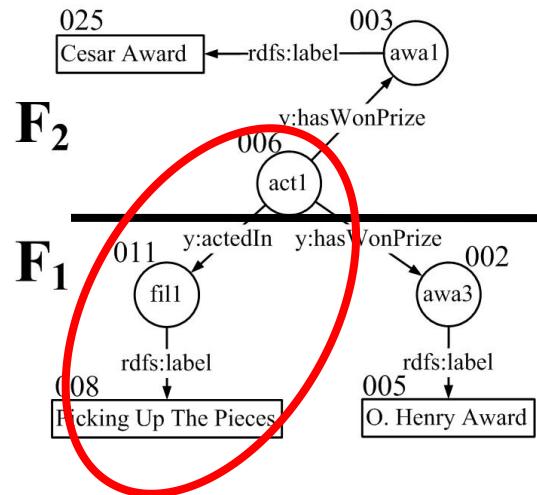
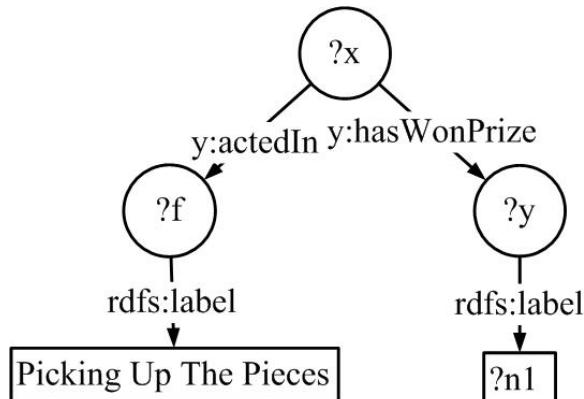
# Definition of Local Partial Match

---

- Given a query  $Q$  with  $n$  vertices  $\{v_1, \dots, v_n\}$  and a subgraph  $PM$  with  $m$  vertices  $\{u_1, \dots, u_m\}$  ( $m \leq n$ ) in a fragment  $F_k$ ,  $PM$  is a local partial match in fragment  $F_k$  if and only if there exists a function  $f : \{v_1, \dots, v_n\} \rightarrow \{u_1, \dots, u_m\} \cup \{\text{NULL}\}$ , where the following conditions hold:
  1. If  $f(v_i)$  is an internal vertex in  $F_k$  and  $v_j$  is a neighbor,  $f(v_j) \neq \text{NULL}$ ;
  2. If  $f(v_i)$  and  $f(v_j)$  are both internal vertices in  $PM$ , then there exists a weakly connected path  $\pi$  between  $v_i$  and  $v_j$  in  $Q$  and each vertex in  $\pi$  maps to an internal vertex of  $F_k$  in  $PM$ .

# Example Local Partial Matches

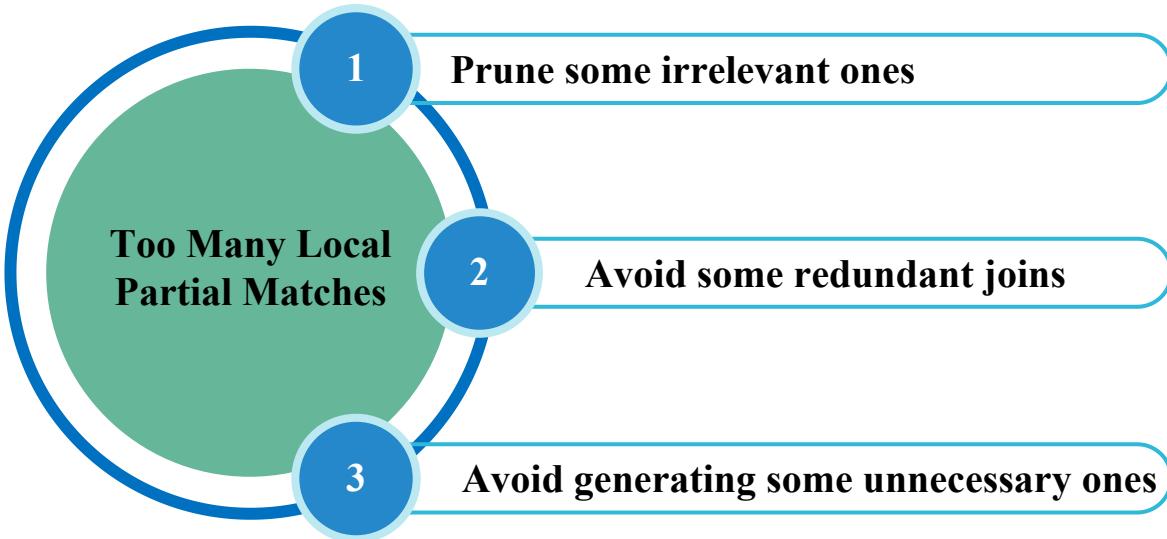
- We propose a distributed RDF system, named gStore<sup>D</sup>, to evaluate SPARQL queries in a “partial evaluation” framework.



There may be too many local partial matches!!!

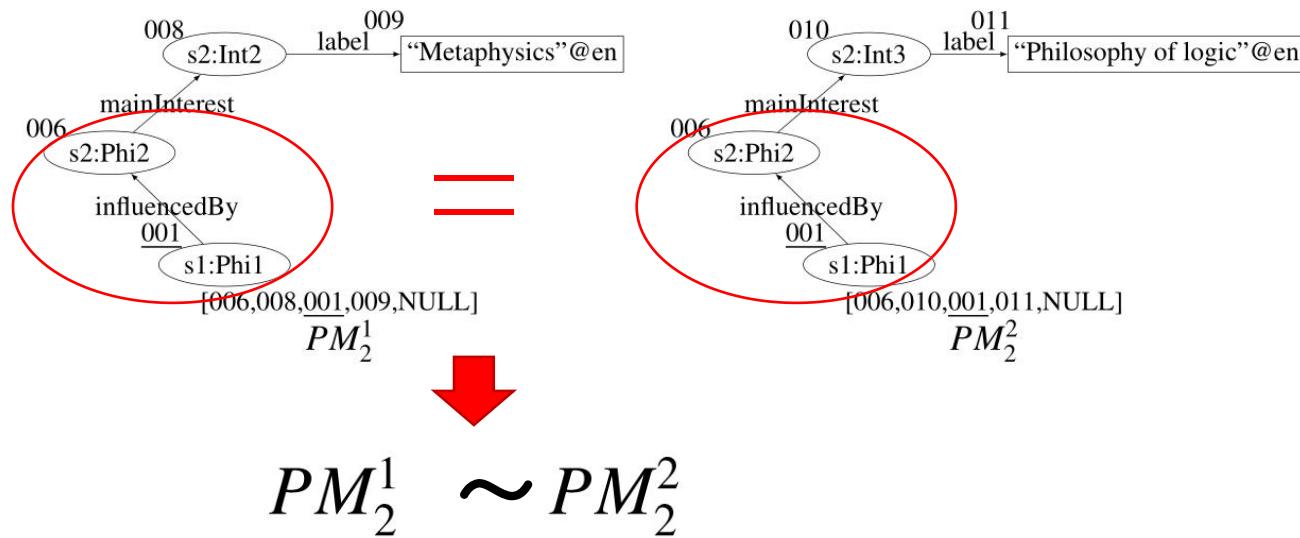
# Three Optimizations

---



# Properties of Local Partial Matches

- Two local partial matches of the **same crossing edges** from the **same fragment** have the same structures.



# Local Partial Match Equivalence Class (LEC)

---

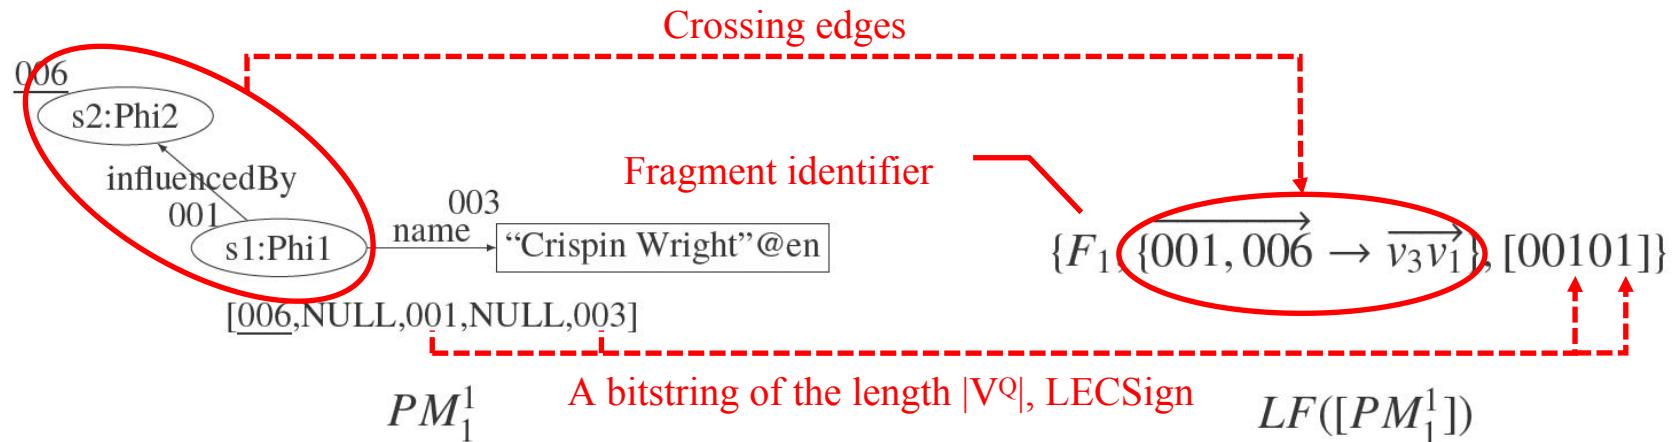
- All local partial matches of the same crossing edges from the same fragment are put into the same class

$$PM_2^1 \sim PM_2^2$$


$$\begin{aligned}[PM_2^1] &= \{PM_2^i \mid PM_2^i \sim PM_2^1\} \\ &= \{PM_2^1, PM_2^2\} \\ &= [PM_2^2]\end{aligned}$$

# LEC Feature

- The structural information of a LEC is maintained in a compact data structure called *LEC feature*



# Joinable

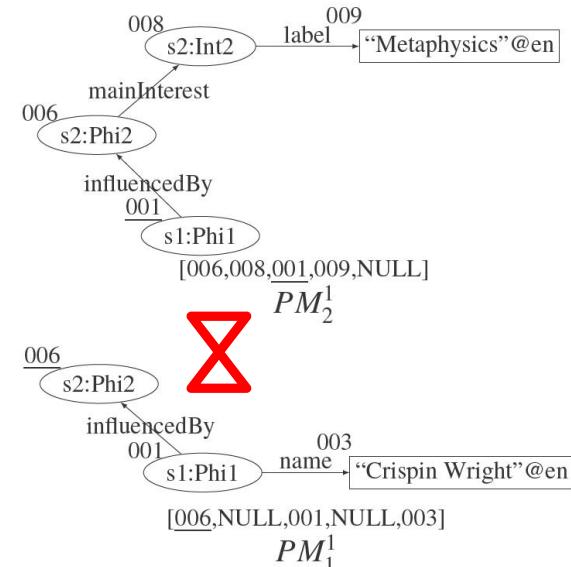
- Two LEC features can join if they share the same crossing edges and all bits in the AND result of their LECSigns are 0

$$LF([PM_2^1]) = \{F_2, \{\overrightarrow{001, 006} \rightarrow \overrightarrow{v_3 v_1}\}, [11010]\}$$

||

$$LF([PM_1^1]) = \{F_1, \{\overrightarrow{001, 006} \rightarrow \overrightarrow{v_3 v_1}\}, [00101]\}$$

$\wedge = [00000]$



# Joinable

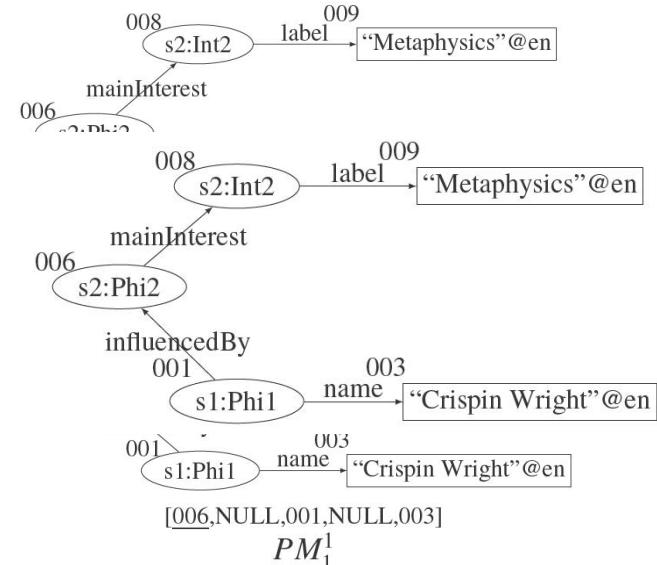
- A local partial match can contribute to a final match if its LEC feature can join with some LEC features and all bits in the OR result of their LECSians are 1

$$LF([PM_2^1]) = \{F_2, \{001, 006 \rightarrow \overrightarrow{v_3 v_1}\}, [11010]\}$$

||

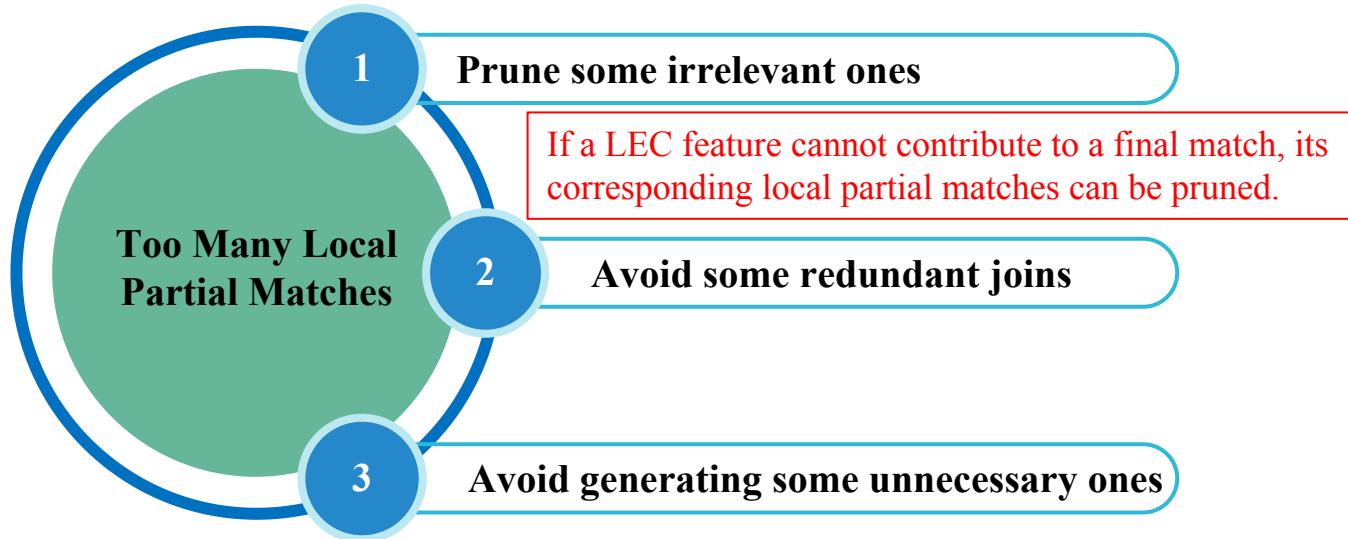
$$LF([PM_1^1]) = \{F_1, \{001, 006 \rightarrow \overrightarrow{v_3 v_1}\}, [00101]\}$$

$v=[11111]$



# Optimizations based on LEC Features

---



The optimization is partition bounded in both response time and data shipment

# LEC Feature-based Local Partial Match Group

---

- If the LECSigns of two local partial matches are the same, they cannot join together

Group 2

$$LF([PM_1^1]) = \{F_1, \overrightarrow{001, 006} \rightarrow \overrightarrow{v_3 v_1}, [00101]\}$$

$$LF([PM_1^2]) = \{F_1, \overrightarrow{001, 012} \rightarrow \overrightarrow{v_3 v_1}, [00101]\}$$

$$LF([PM_1^3]) = \{F_1, \overrightarrow{006, 005} \rightarrow \overrightarrow{v_1 v_2}, [01010]\}$$

Group 3

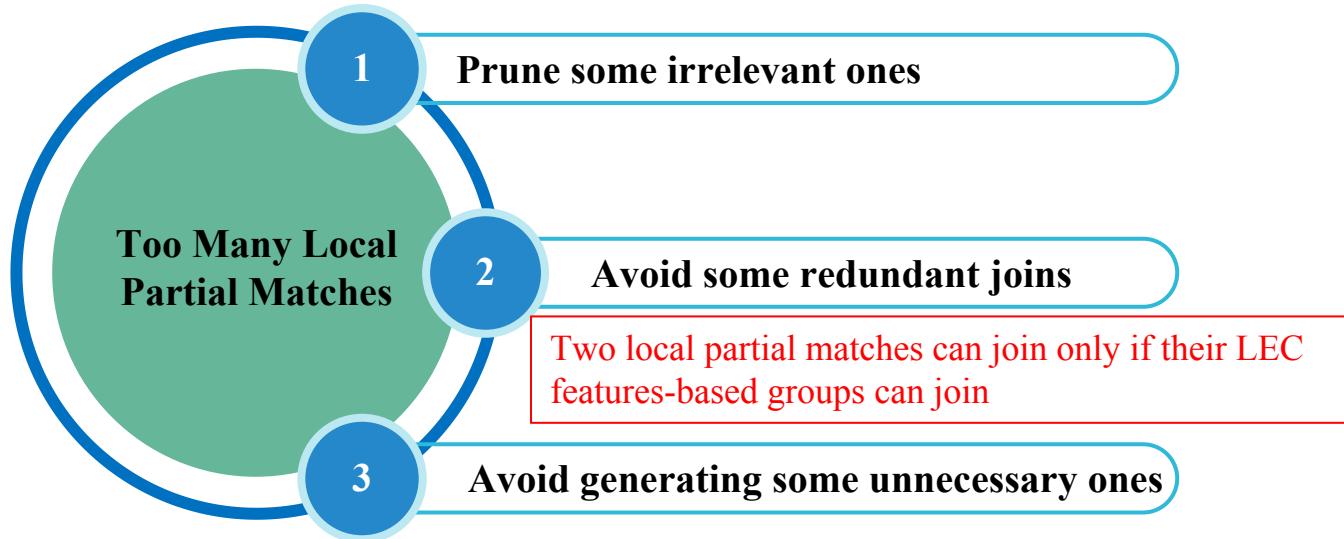
$$LF([PM_2^1]) = \{F_2, \overrightarrow{001, 006} \rightarrow \overrightarrow{v_3 v_1}, [11010]\}$$

$$LF([PM_3^1]) = \{F_3, \overrightarrow{001, 012} \rightarrow \overrightarrow{v_3 v_1}, [11010]\}$$

$$LF([PM_3^2]) = \{F_3, \overrightarrow{014, 013} \rightarrow \overrightarrow{v_1 v_2}, [01010]\}$$

# Optimizations based on LEC Feature-based Group

---



# Assembling Internal Candidates

---

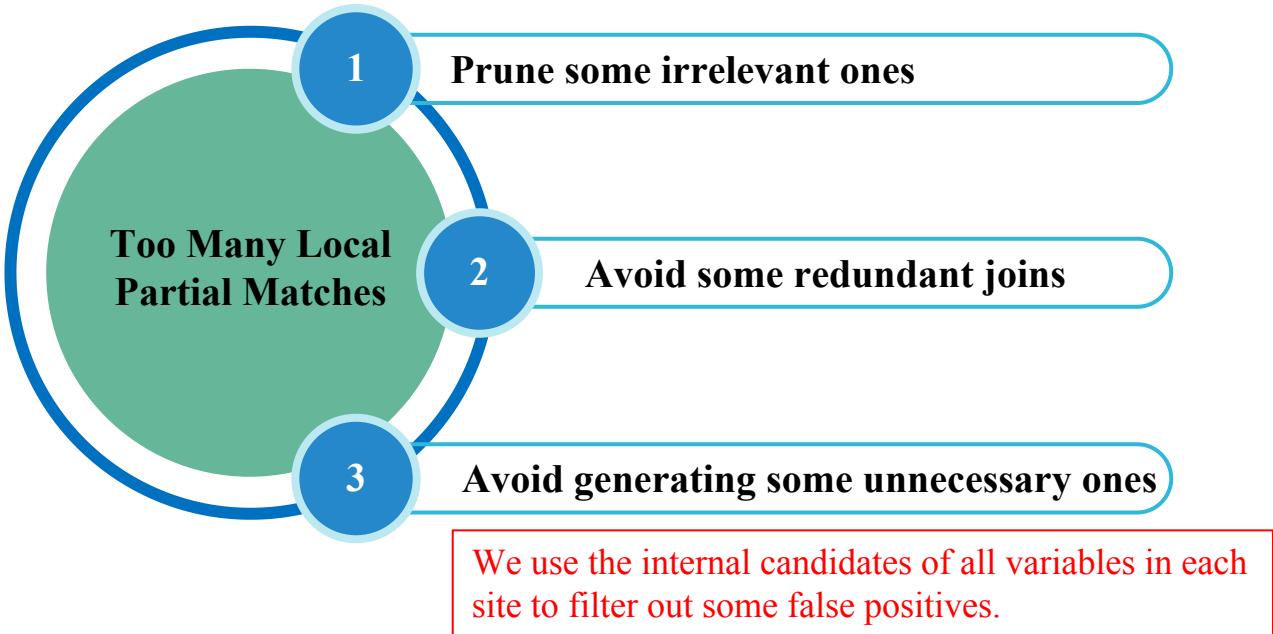
A vertex occurring in a final match must be an internal vertex of a fragment

We can prune a candidate of a variable that is only an extended vertex of fragments

We can avoid generating the local partial matches from the above candidates

# Optimizations based on Internal Candidates

---



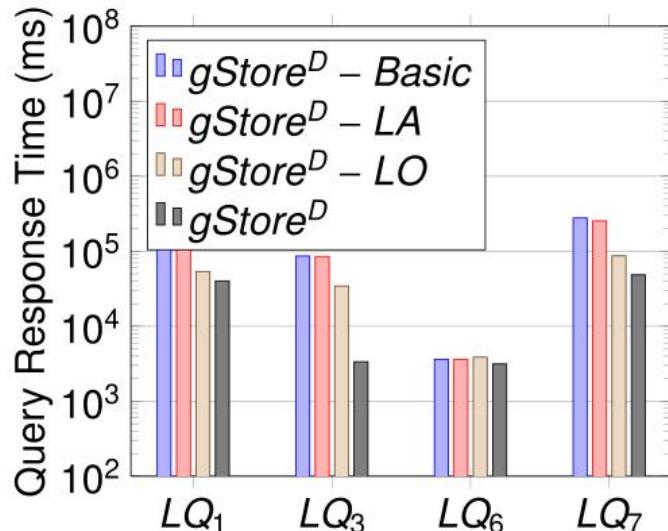
# Experiments

---

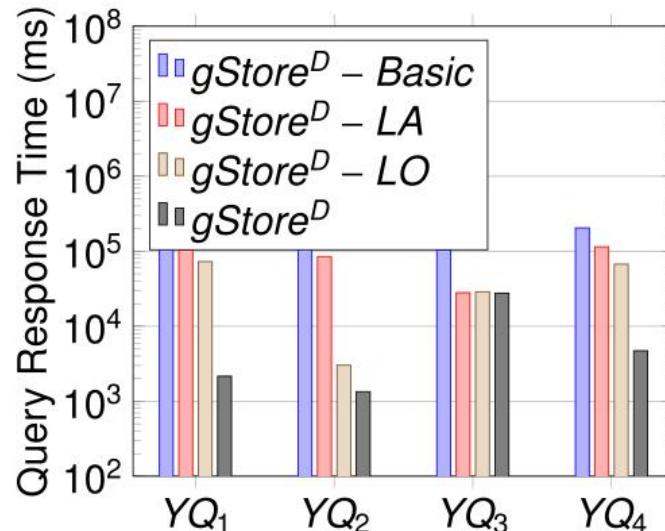
- **Datasets:** YAGO2 (about 284 millions triples), BTC (about 1 billion triples) and LUBM (100 millions to 1 billion triples)
- **Competitor:** DREAM, S2X, S2RDF and CliqueSquare
- **Environment:** 12machines running Linux in a LAN

# Evaluation of Different Optimizations

---

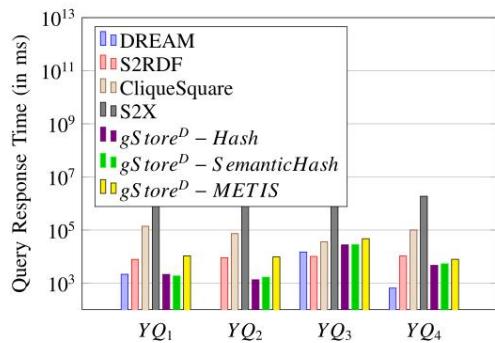


(a) LUBM 100M

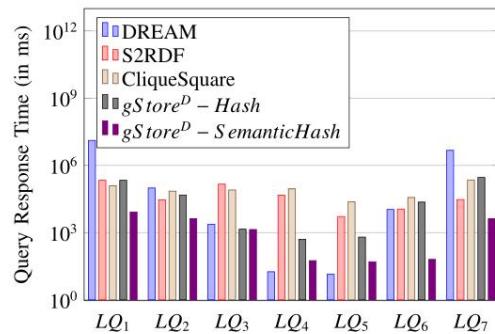


(b) YAGO2

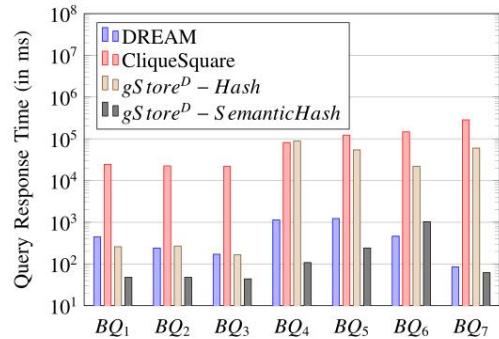
# Online Performance Comparison



(a) YAGO2



(b) LUBM 1B



(c) BTC

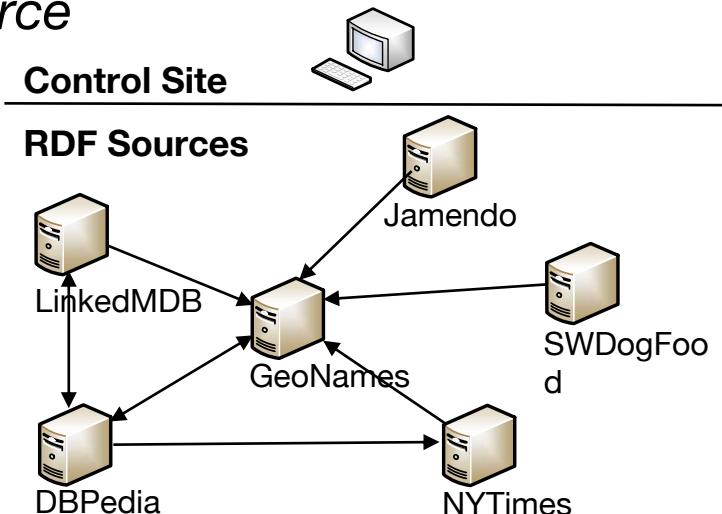
# **Part 3**

Federated Systems

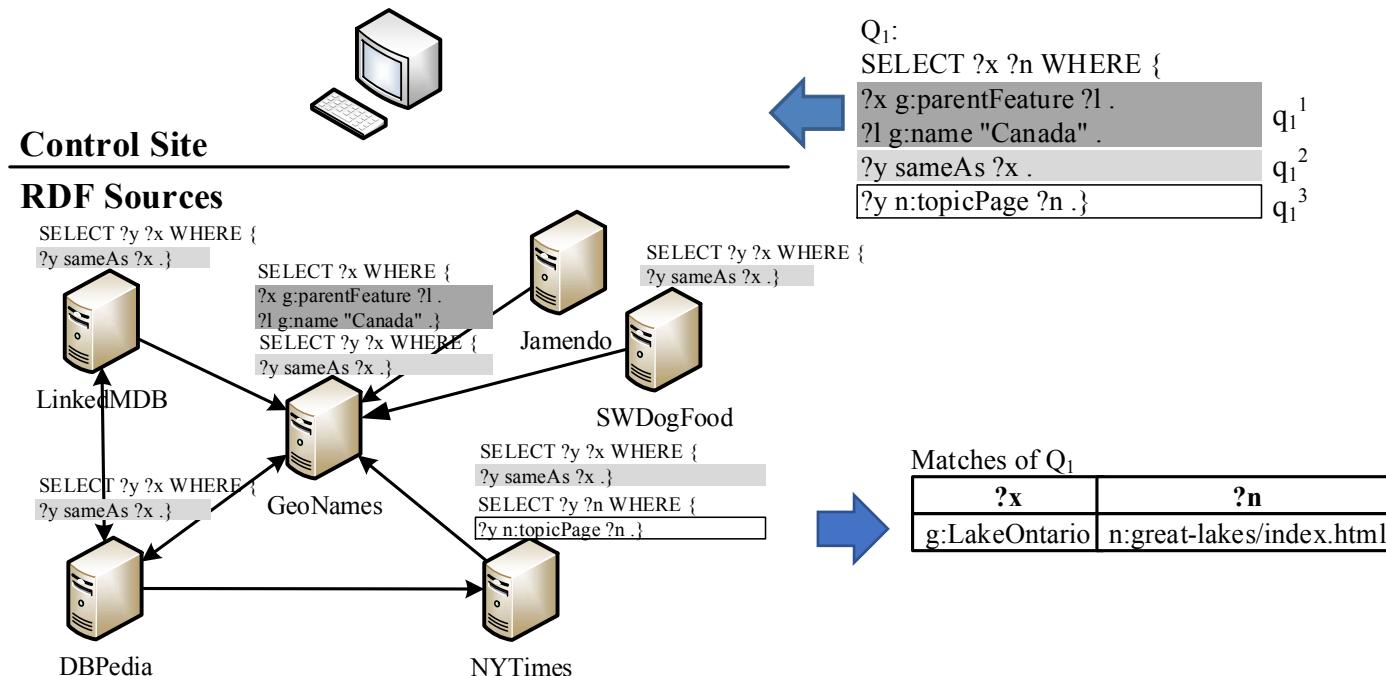
# Federated RDF Systems

---

- A federated RDF system consists of multiple "autonomous" sites which store their own RDF data and provide SPARQL query interfaces. Here, an autonomous site with a SPARQL interface is called an *RDF source*



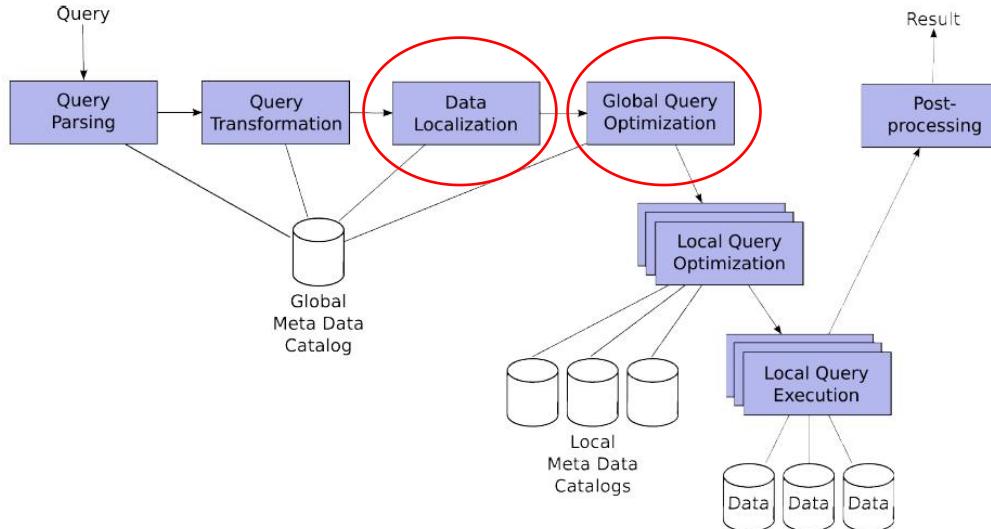
# Federated SPARQL Query Processing



# Existing Approaches

---

- ❑ Most of them focus on query decomposition and source selection



# Federated RDF Systems

---

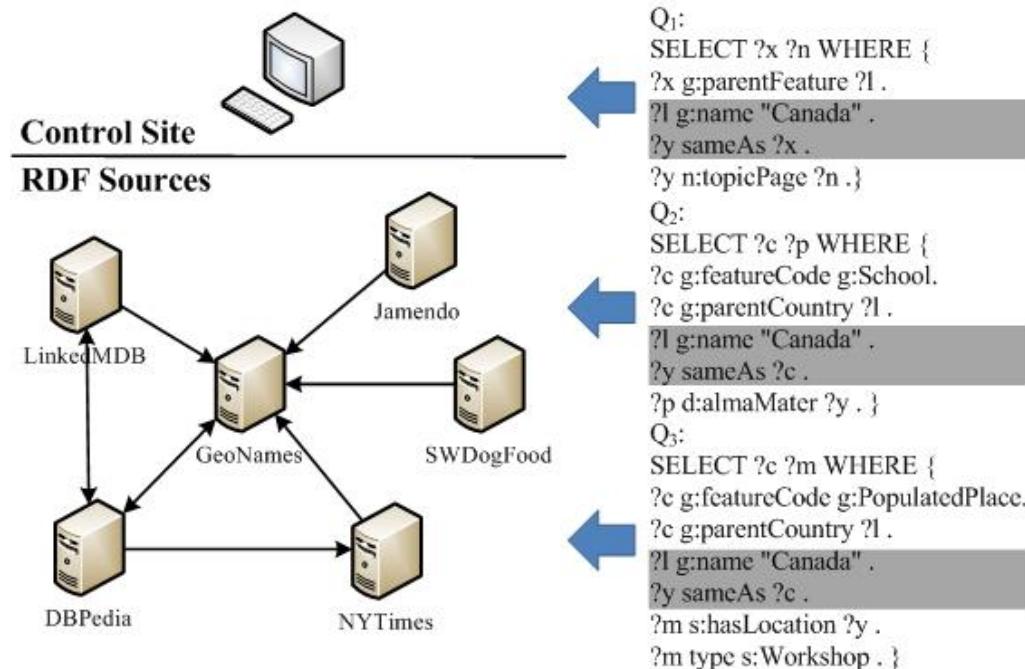
## Metadata-based approaches

- Use the metadata to determine which sources are relevant
  - DARQ [Quilizt & Leser, ESWC 2008]
  - QTree [Harth et al., WWW 2010; Prasser et al., EDBT 2012]
  - ...

## ASK query-based approach

- Asking whether or not a triple pattern has an answer at a source
  - FedX [Schwarte et al., ISWC 2011]

# Multi-Query Optimization [Peng et al., DASFAA 2018, TKDE 2021]



The shaded triple patterns correspond to the common subgraph during each iteration.

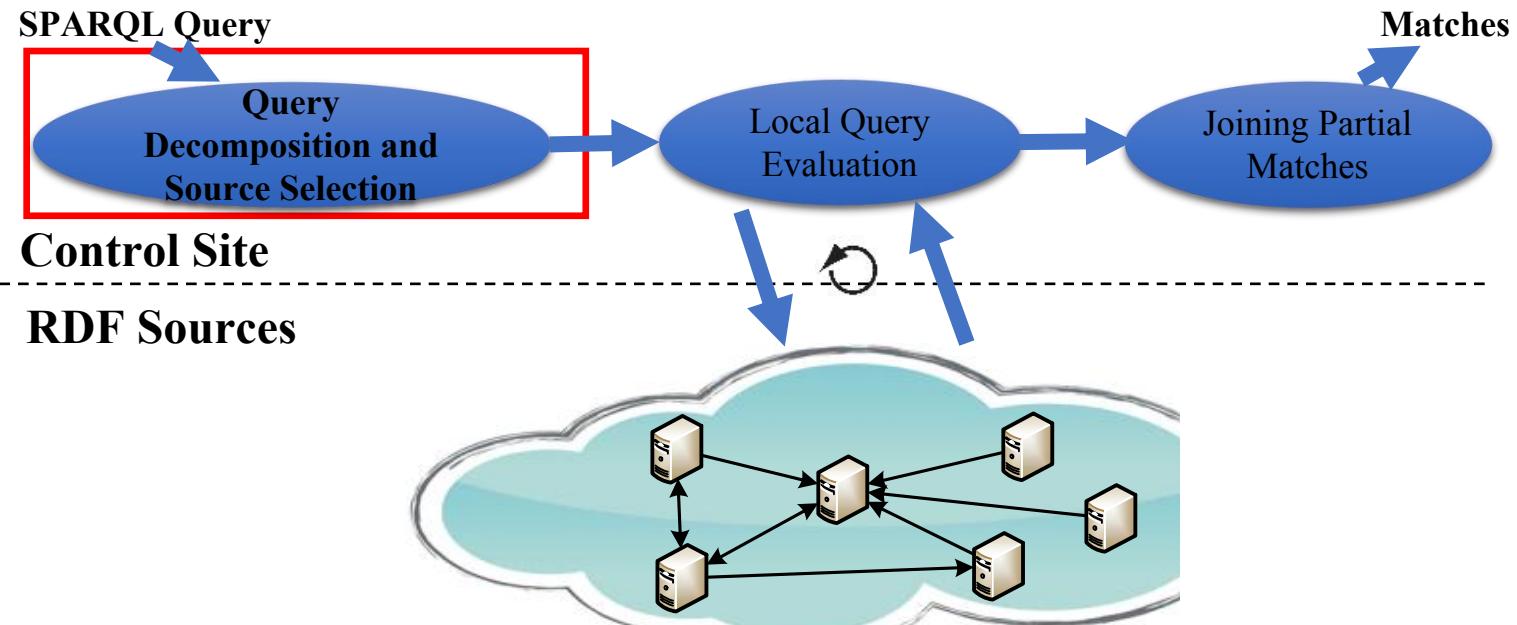
# Main Ideas

---

- A query rewriting-based approach for query evaluation on RDF sources while considering the cost of both query evaluation and data shipment
- An effective method to use the interconnection topology between RDF sources to filter out irrelevant sources
- An efficient method to share the common computation of intermediate results joining.

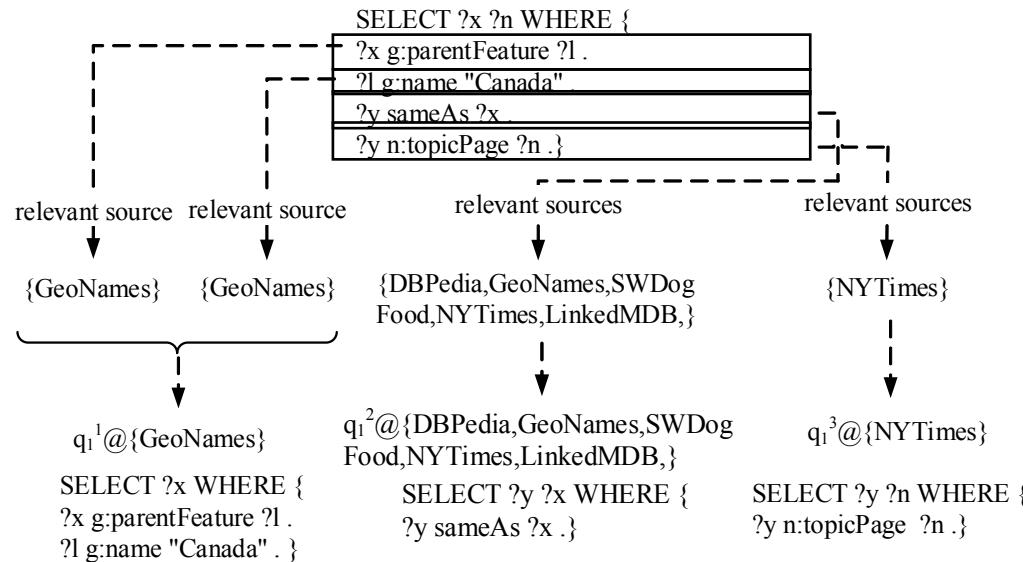
# Framework

---

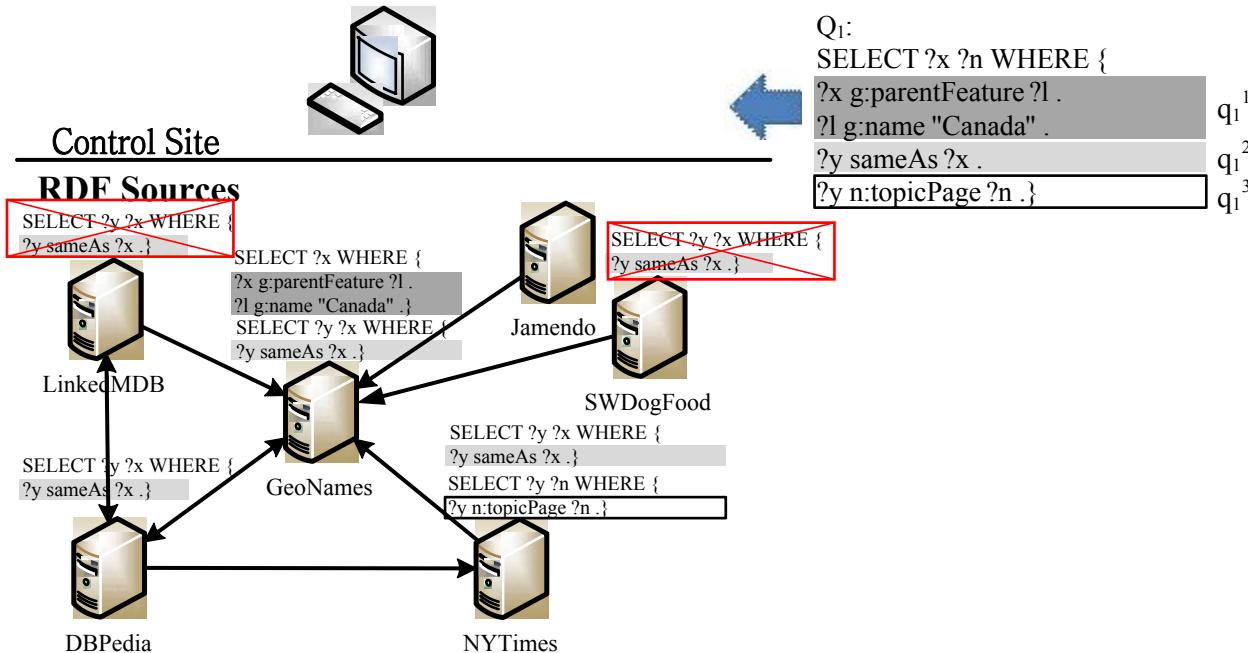


# Basic Query Decomposition and Source Selection

- Most existing solutions decompose the input query based on the triple patterns

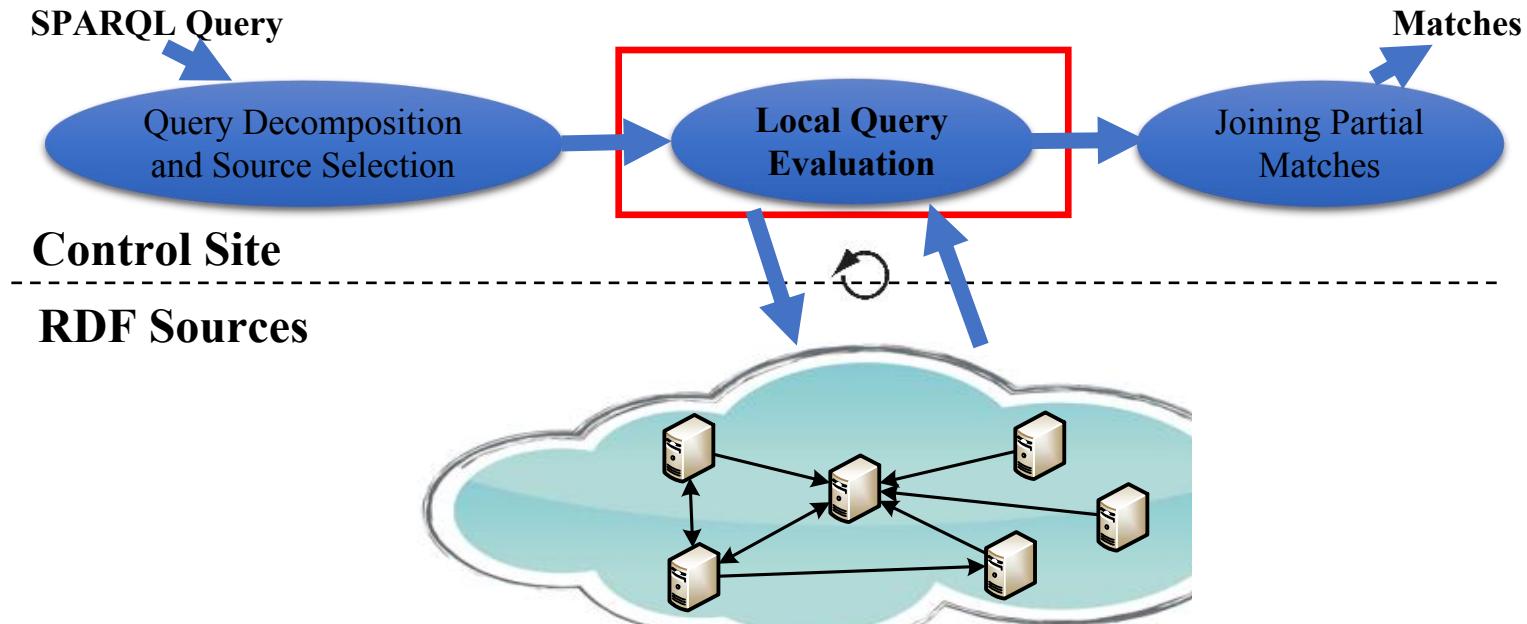


# Source Topology-based Query Decomposition and Source Selection



# Framework

---



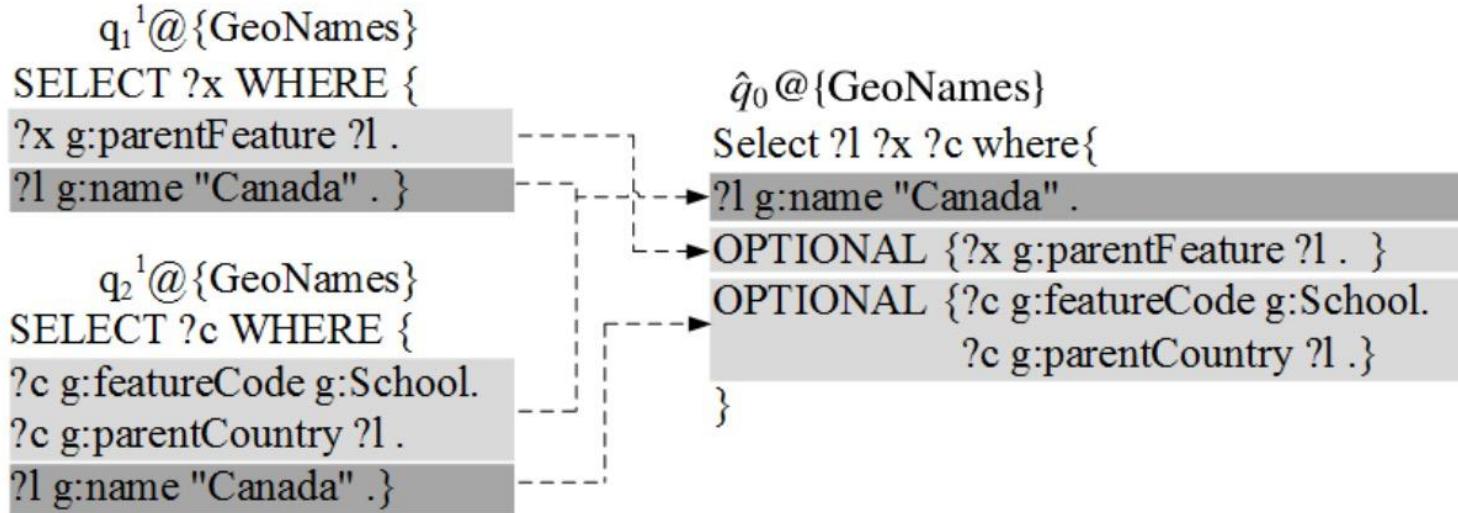
# Local Query Evaluation

---

- When multiple SPARQL queries are posed simultaneously, there is room for sharing computation when executing these queries
- We proposes a **cost-driven query rewriting scheme** to rewrite them (at the control site) into fewer SPARQL queries, which can reduce the number of remote requests and improve the performances

# OPTIONAL-Based Rewriting

---



# OPTIONAL-UNION-Based Rewriting

---

$q_1^1 @ \{\text{GeoNames}\}$

```
SELECT ?x WHERE {  
?x g:parentFeature ?l .  
?l g:name "Canada" . }
```

$q_2^1 @ \{\text{GeoNames}\}$

```
SELECT ?c WHERE {  
?c g:featureCode g:School.  
?c g:parentCountry ?l .  
?l g:name "Canada" . }
```

$\hat{q}_0 @ \{\text{GeoNames}\}$

```
Select ?l ?x ?c where {
```

```
?l g:name "Canada" .
```

```
OPTIONAL {
```

```
{?x g:parentFeature ?l . }
```

```
UNION
```

```
{?c g:featureCode g:School.?c g:parentCountry ?l . }}
```

```
}
```

# FILTER-based Rewriting

$q_2^1 @ \{\text{GeoNames}\}$

```
SELECT ?c WHERE {
  ?c g:featureCode [g:School .]
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .}
```

$q_3^1 @ \{\text{GeoNames}\}$

```
SELECT ?c WHERE {
  ?c g:featureCode [g:PopulatedPlace .]
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .}
```

$\hat{q}_1 @ \{\text{GeoNames}\}$

```
Select ?l ?c ?f where{
  ?c g:featureCode ?f .
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .
  Filter(?f = [g:School] || ?f = [g:PopulatedPlace])}
```

# VALUES-based Rewriting

---

$q_2^1 @ \{\text{GeoNames}\}$

```
SELECT ?c WHERE {
  ?c g:featureCode [g:School] .
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .}
```

$q_3^1 @ \{\text{GeoNames}\}$

```
SELECT ?c WHERE {
  ?c g:featureCode [g:PopulatedPlace] .
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .}
```

$\hat{q}_1 @ \{\text{GeoNames}\}$

```
Select ?l ?c ?f where {
  ?c g:featureCode ?f .
  ?c g:parentCountry ?l .
  ?l g:name "Canada" .}
```

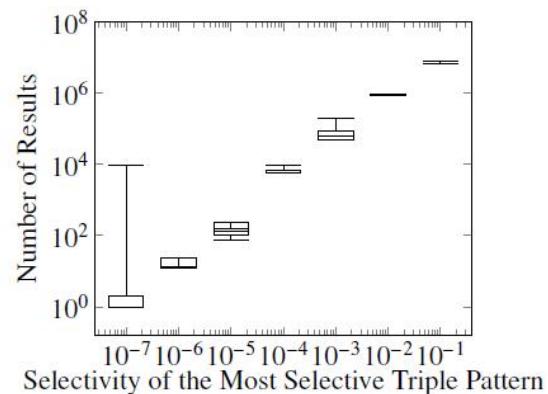
```
VALUES (?f) { (g:School) (g:PopulatedPlace) }
```

# Cost Model for BGPs

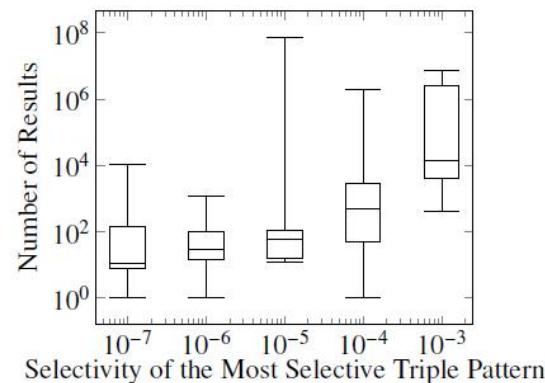
- The cardinality of evaluating a basic graph pattern  $Q$  is as follows.

$$card(Q) = \min_{e \in E(Q)} \{sel(e)\}$$

where  $sel(e)$  is the selectivity of triple pattern  $e$  in  $Q$ .



(a) DBpedia



(b) WatDiv 10M

**Fig. 6.** Relationship Between the Cardinality and the Most Selective Triple Pattern

# Cost Model for Query Rewriting

---

- Hence, given a SPARQL  $Q$ , its cost  $\text{cost}_{LE}(Q)$  for local evaluation is defined as follows

$$\text{card}(Q) = \begin{cases} \min_{e \in E(Q)} \{ \text{sel}(e) \} & \text{if } Q \text{ is a BGP;} \\ \min \{ \text{card}(Q_1), \text{card}(Q_2) \} & \text{if } Q = Q_1 \text{ AND } Q_2; \\ \text{card}(Q_1) + \text{card}(Q_2) & \text{if } Q = Q_1 \text{ UNION } Q_2; \\ \text{card}(Q_1) + \Delta_1 & \text{if } Q = Q_1 \text{ OPT } Q_2; \\ \text{card}(Q_1) + \Delta_2 & \text{if } Q = Q_1 \text{ FILTER } F; \\ \min \{ \text{card}(Q), |D| \} & \text{if } Q = \text{VALUES } \overrightarrow{W} D; \end{cases}$$

where  $T_{CPU}$  is the CPU unit time to construct a result, and  $\Delta_3$  and  $\Delta_4$  are empirically trivial values

# Cost Model for Data Shipment

---

- Based on the above, the cost of data shipment for a given query is as follows

$$cost_{DS}(\hat{q}) = \text{card}(\hat{q}) \times T_{MSG} = \min_{e \in p} \{sel(e)\} \times T_{MSG}$$

where  $T_{MSG}$  is the unit time to transmit a data unit.

# Cost Model for Query Rewriting

---

- Given a set of subqueries on a RDF source, if  $p$  is their common subgraph among these queries, we rewrite them into a SPARQL query. The cost of the rewriting is the cost of the rewritten query is as follows.

$$\begin{aligned}cost(\hat{q}) &= cost_{LE}(\hat{q}) + cost_{DS}(\hat{q}) \\&= \min_{e \in p} \{sel(e)\} \times (T_{CPU} + T_{MSG}) = cost(p).\end{aligned}$$

# Query Rewriting

---

- Given a set of  $n$  subqueries, the objective is to find the set of  $m$  rewritten queries ( $m \leq n$ ) with the smallest rewriting cost.
- We can prove that finding the optimal set of rewritten queries is NP-complete, and propose a greedy heuristic algorithm to find an approximately optimal solution

# Query Rewriting Example

$q_1^1 @ \{ \text{GeoNames} \}$   
SELECT ?x WHERE {  
?x g:parentFeature ?l .  
?l g:name "Canada" . }

$q_2^1 @ \{ \text{GeoNames} \}$   
SELECT ?c WHERE {  
?c g:featureCode g:School.  
?c g:parentCountry ?l .  
?l g:name "Canada" . }

$q_3^1 @ \{ \text{GeoNames} \}$   
SELECT ?c WHERE {  
?c g:featureCode g:PopulatedPlace.  
?c g:parentCountry ?l .  
?l g:name "Canada" . }

## First equivalence class

$q_1^1 @ \{ \text{GeoNames} \}$   
SELECT ?x WHERE {  
?x g:parentFeature ?l .  
?l g:name "Canada" . }

$q_1^1 @ \{ \text{GeoNames} \}$   
SELECT ?x WHERE {  
?x g:parentFeature ?l .  
?l g:name "Canada" . }

## Second equivalence class

$q_2^1 @ \{ \text{GeoNames} \}$   
SELECT ?c WHERE {  
?c g:featureCode g:School.  
?c g:parentCountry ?l .  
?l g:name "Canada" . }

$\hat{q}_1 @ \{ \text{GeoNames} \}$   
Select ?l ?x ?c ?f where {  
?l g:name "Canada".  
?c g:featureCode ?f .  
?c g:parentCountry ?l .  
?l g:name "Canada" . }

$q_3^1 @ \{ \text{GeoNames} \}$   
SELECT ?c WHERE {  
?c g:featureCode g:PopulatedPlace.  
?c g:parentCountry ?l .  
?l g:name "Canada" . }

$\hat{q}_1 @ \{ \text{GeoNames} \}$   
Select ?l ?c ?f where {  
?l g:name "Canada".  
?c g:featureCode ?f .  
?c g:parentCountry ?l .  
VALUES (?f) {(g:School) (g:PopulatedPlace)}  
}

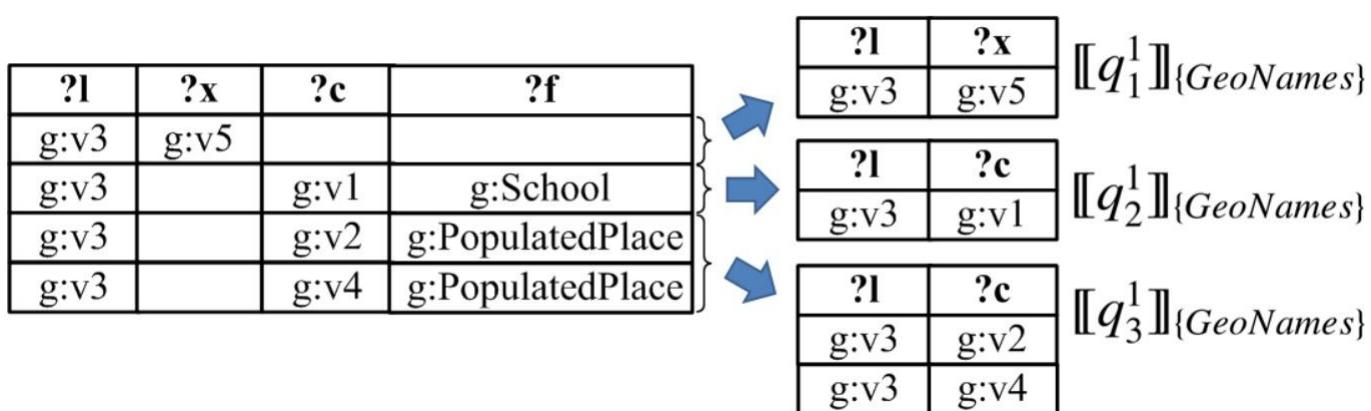
$\hat{q}_2 @ \{ \text{GeoNames} \}$

Select ?l ?x ?c ?f where {  
?l g:name "Canada" .  
OPTIONAL { {?x g:parentFeature ?l . } UNION  
{?c g:featureCode ?f .  
?c g:parentCountry ?l .  
VALUES (?f) {(g:School) (g:PopulatedPlace)} } } }



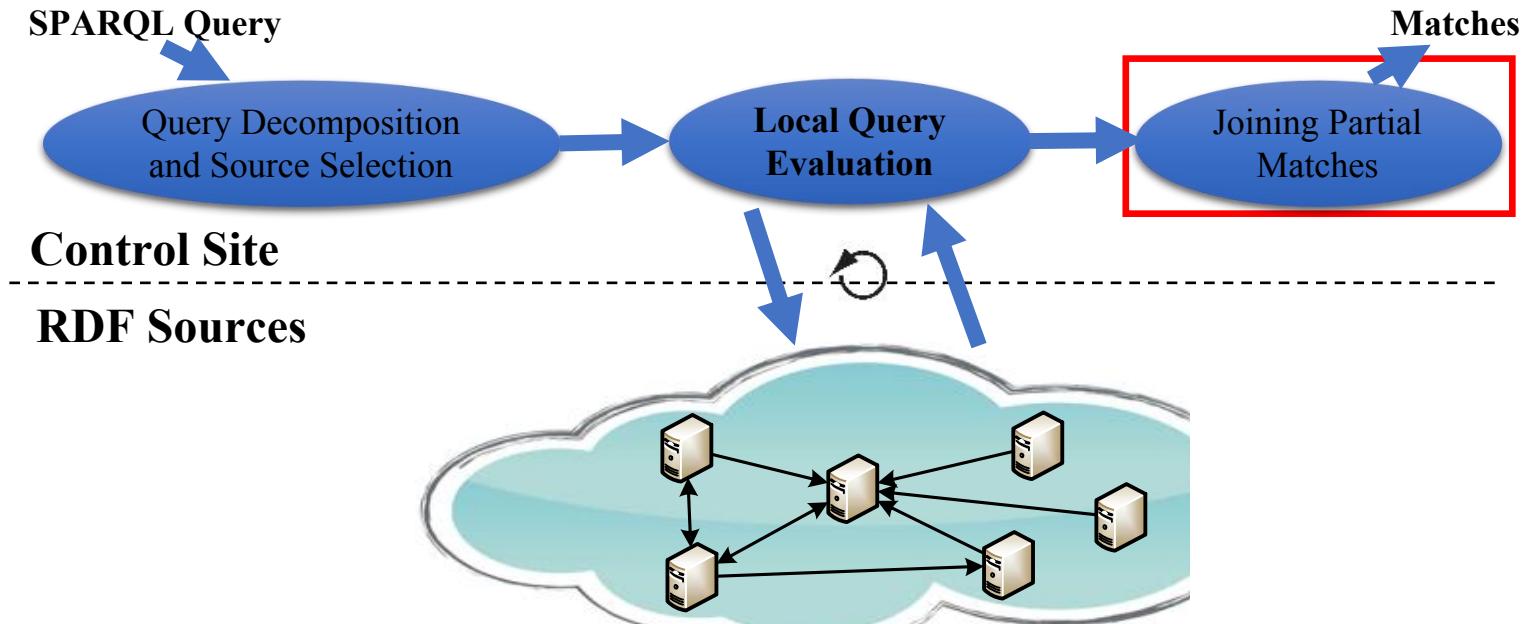
# Postprocessing

---



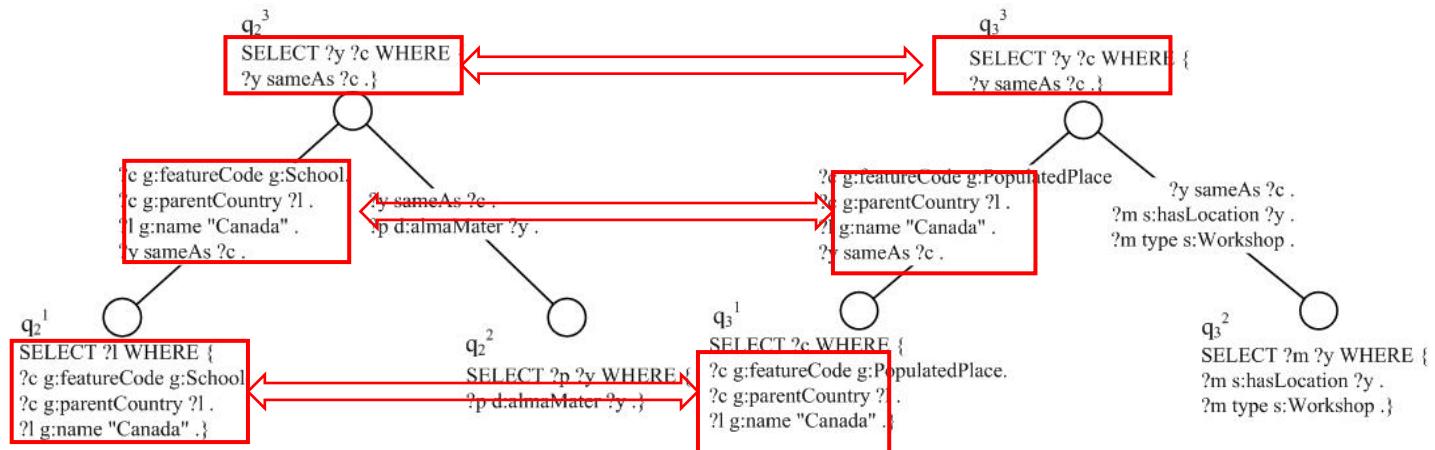
# Framework

---



# Joining Partial Matches

- There may exist some **common computation** in joining partial matches of different queries



# Joining Partial Matches

---

- Formally,

$$\begin{array}{c} \llbracket q_2^1 \rrbracket \bowtie \llbracket q_2^3 \rrbracket \\ \qquad \qquad \qquad \left.\right\downarrow \\ \llbracket q_3^1 \rrbracket \bowtie \llbracket q_3^3 \rrbracket \end{array} \xrightarrow{\hspace{1cm}} (\llbracket q_2^1 \rrbracket \cup \llbracket q_3^1 \rrbracket) \bowtie (\llbracket q_2^3 \rrbracket \cup \llbracket q_3^3 \rrbracket)$$

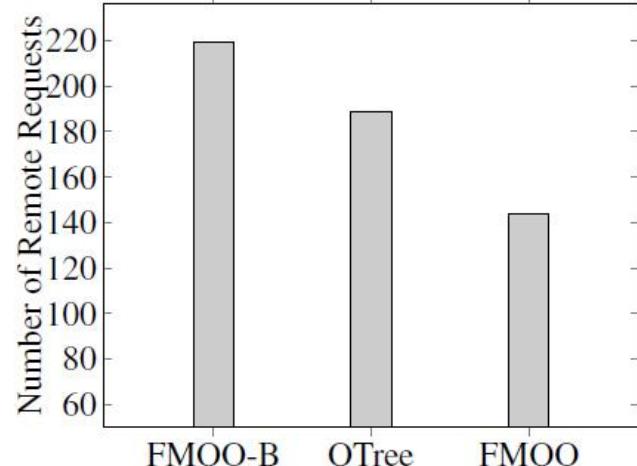
# Experiments

---

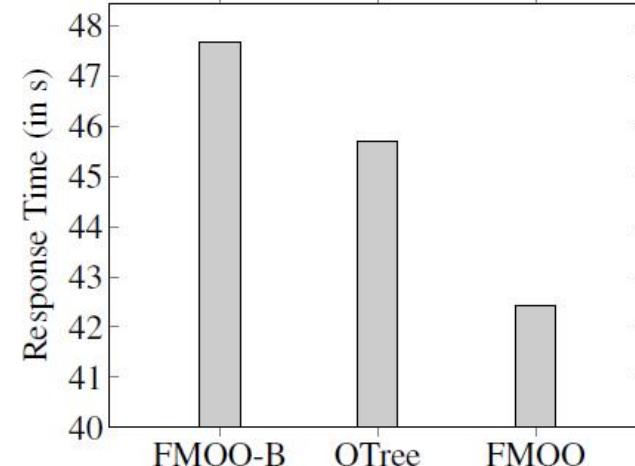
- **Datasets:** WatDiv, FedBench and LargeRDFBench
- **Competitor:** FedX, SPLENDID and HiBISCuS
- **Environment:** a cluster of machines running Linux, each of which installs Sesame 2.7 to build up an RDF source.

# Effect of the Query Decomposition and Source Selection Technique

---



(a) Number of Remote Requests

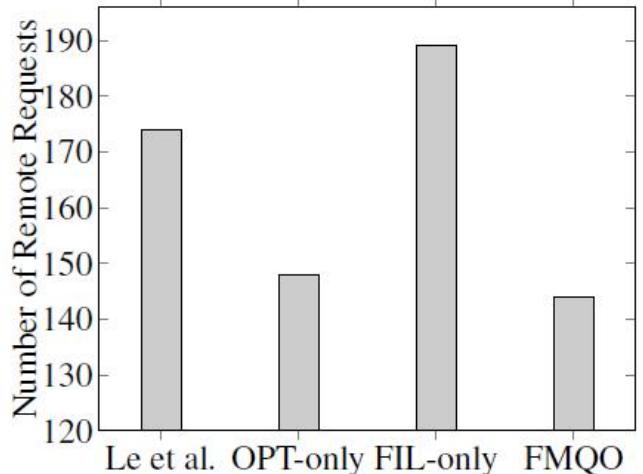


(b) Response Time

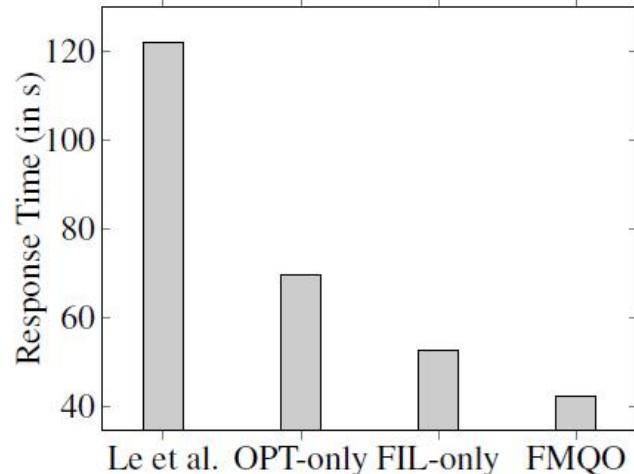
**Fig. 8.** Evaluating Source Topology-based Source Selection Technique

# Effect of the Rewriting Strategies

---



(a) Number of Remote Requests

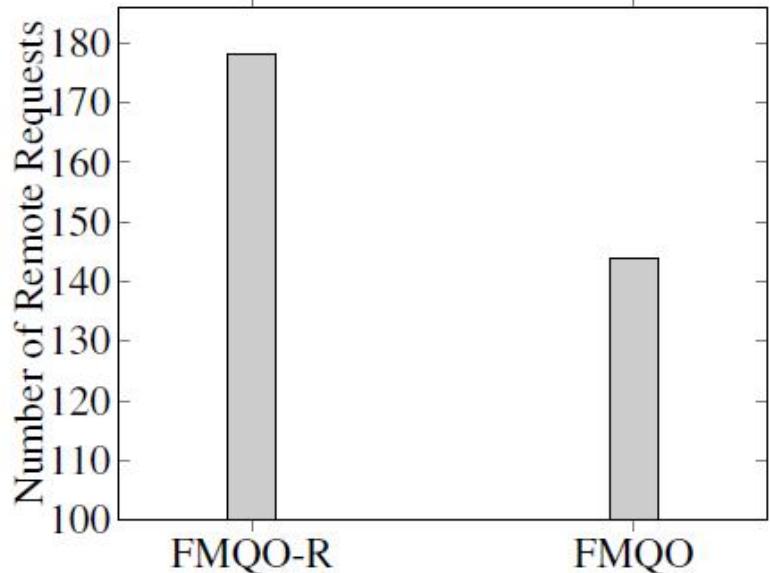


(b) Response Time

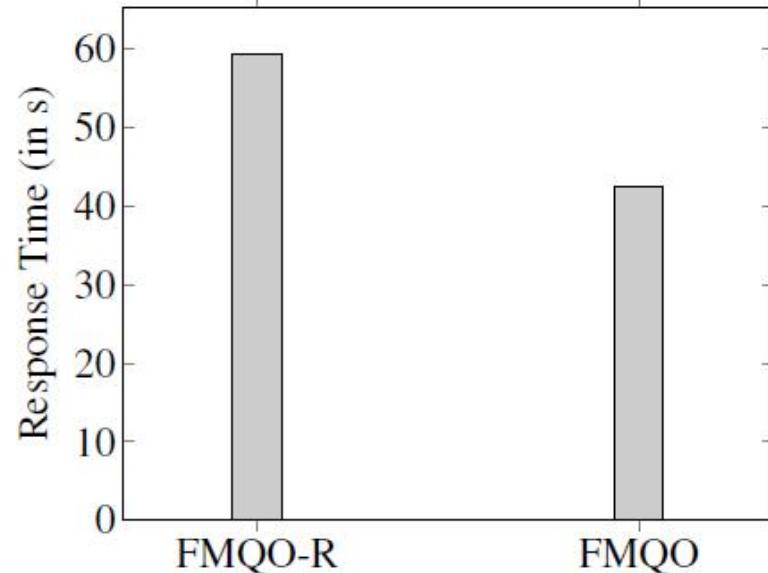
**Fig. 9.** Evaluating Different Rewriting Strategies

# Evaluation of the Cost Model

---



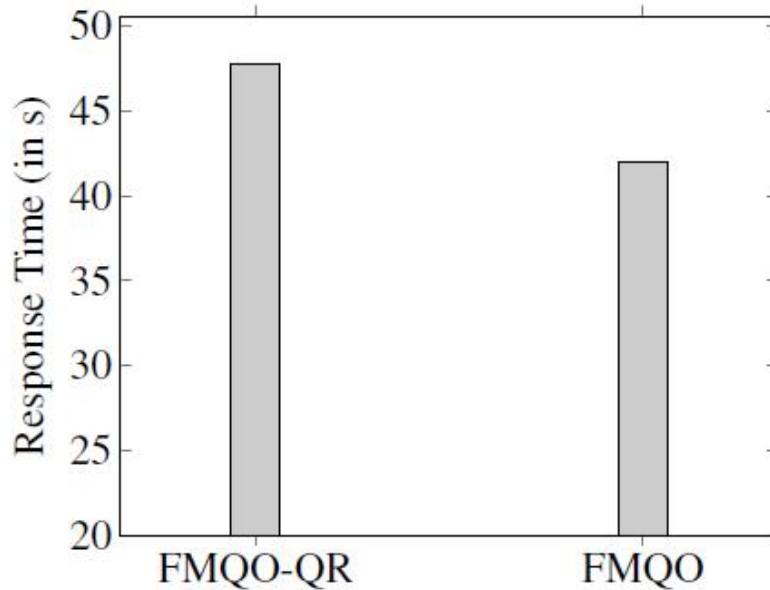
(a) Number of Remote Requests



(b) Response Time

# Effect of Optimization Techniques for Joins

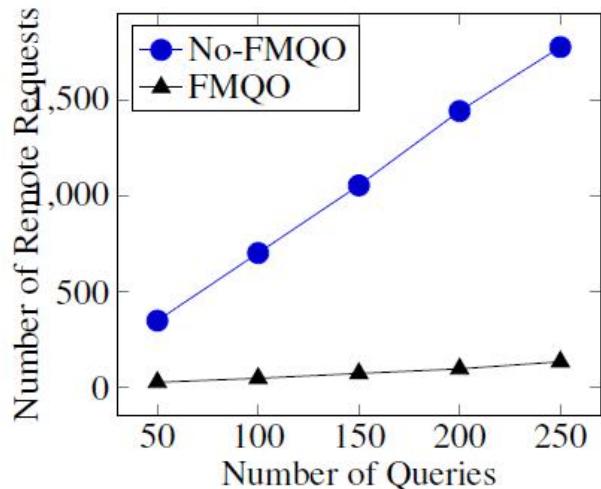
---



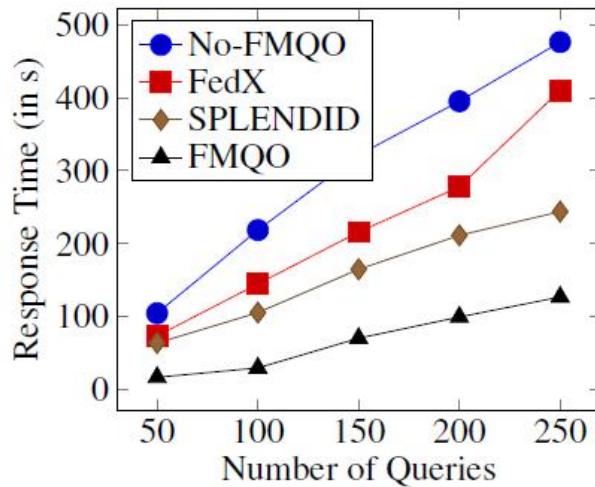
**Fig. 11.** Effect of Optimization Techniques for Joins

# FedBench (Cross Domain)

---



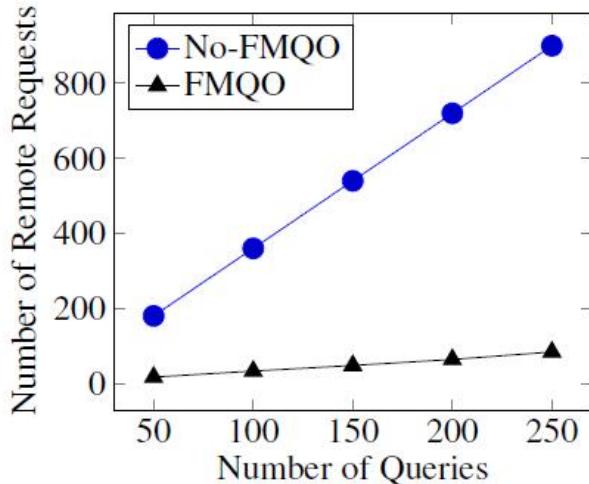
(a) Number of Remote Requests



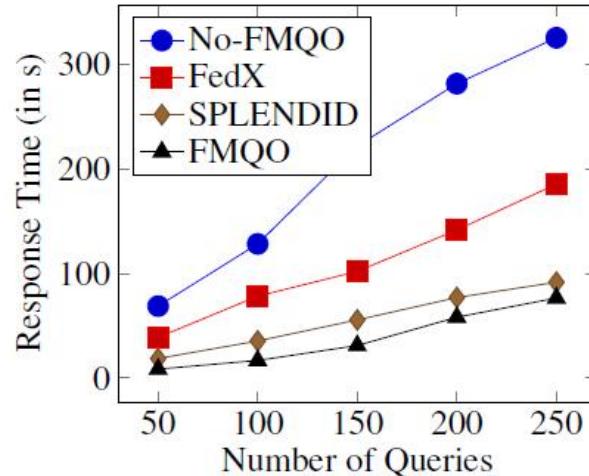
(b) Response Time

# FedBench (Life Science)

---



(a) Number of Remote Requests



(b) Response Time

# **Part 4**

## Conclusions

# Conclusions

---

- We introduce **partitioning-agnostic** and **federated** RDF systems and discuss how to optimize query processing

# Future Work

---

- There are many distributed database management system architectures which are further used for RDF graphs
  - 1. Parallel Computing
  - 2. Edge Computing

.....

# THANK YOU

---

# References

---

- Peng Peng, Lei Zou, Runyu Guan. Accelerating Partial Evaluation in Distributed SPARQL Query Evaluation. ICDE 2019: 112-123
- Peng Peng, Lei Zou, M. Tamer Özsü, Lei Chen, Dongyan Zhao. Processing SPARQL queries over distributed RDF graphs. VLDB J. 25(2): 243-268 (2016)
- Bastian Quilitz, Ulf Leser. Querying Distributed RDF Data Sources with SPARQL. ESWC 2008: 524-538
- Andreas Harth, Katja Hose, Marcel Karnstedt, Axel Polleres, Kai-Uwe Sattler, Jürgen Umbrich. Data summaries for on-demand queries over linked data. WWW 2010: 411-420
- Fabian Prasser, Alfons Kemper, Klaus A. Kuhn. Efficient distributed query processing for autonomous RDF databases. EDBT 2012: 372-383
- Andreas Schwarte, Peter Haase, Katja Hose, Ralf Schenkel, Michael Schmidt. FedEx: Optimization Techniques for Federated Query Processing on Linked Data. ISWC (1) 2011: 601-616
- Peng Peng, Qi Ge, Lei Zou, M. Tamer Özsü, Zhiwei Xu, Dongyan Zhao. Optimizing Multi-Query Evaluation in Federated RDF Systems. IEEE Trans. Knowl. Data Eng. 33(4): 1692-1707 (2021)
- Peng Peng, Lei Zou, M. Tamer Özsü, Dongyan Zhao. Multi-query Optimization in Federated RDF Systems. DASFAA (1) 2018: 745-765